

The labelling workbench — operating guide

Accuracy Lab · how to use it · September 2026

1. What you are looking at

The workbench does one thing: **judge, one at a time, the findings the engine produces on real pages**. Your judgement becomes the ground truth we measure precision against, and the cases you decide become automated tests.

Three concepts are enough to work:

- **Snapshot** — a real page, *frozen*: HTML, CSS, images and response headers saved at capture time. It is not a link to the live site: it is the same page, bit for bit, every time you reopen it. That is what makes a judgement repeatable.
- **Finding** — a single claim by the engine about a single element of that snapshot: WCAG criterion, technique, selector, evidence.
- **Label** — your verdict on the finding: **tp** (the finding is correct), **fp** (it is wrong), **uncertain** (the evidence is not enough to decide).

You are not assessing the site: you are assessing *the finding*. The question is always and only "is the engine right to say this, about this element?".

2. Access and conditions

- **Address:** <https://lab.accessibit.com>
- **Credentials:** we send them to you separately from this document. **Each expert has their own account:** labels are signed with your name, which exists to distinguish who judged what — not to assess you.

Use Chrome or a recent Chromium browser: frozen pages are rendered in an iframe, and Chromium is the engine our entire ground truth is defined against. Nothing needs installing, and no network access to the original sites is required — the pages are already in our hands.

You work in isolation, deliberately. You do not see the other experts' judgements, and the counter at the top right counts *your* labels. Knowing that a colleague has already decided a case — or which way — anchors a judgement just as effectively as seeing the answer before answering. If two of you disagree, we genuinely want to know: see section 6.

Do not use AI assistants to form your judgements.

The findings in this corpus **have already been analysed by language models**, and their opinions are stored separately, excluded from every metric and **deliberately not shown to you** — for the same reason you do not see your colleagues' judgements: "the machine says fp" anchors a decision too. What we are buying is the one thing those models cannot produce: **the judgement of a human expert**. A label formed with the help of an AI adds nothing to work already done — it duplicates it, and in doing so it silently corrupts the ground truth everything else is measured against, because it enters the system carrying the authority of a human opinion without being one.

This applies to the verdict and to the written reason alike. You are of course free to consult the normative documentation, the W3C specifications and your usual working tools, screen readers included.

If we detect patterns consistent with automated generation, the contract is terminated.

3. The interface

The screenshot displays the Accuracy Lab interface. On the left, a queue of findings is listed, each with a site name, criterion, and technique. The main area shows a frozen page with a highlighted element. The right panel provides detailed information about the finding, including the key, code, type, selector, message, outerHTML, attributes, and computed styles. A cookie notice is also visible at the bottom of the main area.

The working screen. On the left, the queue of findings; in the centre, the frozen page with the reported element outlined in red; on the right, the full evidence and the verdict buttons. The two vertical separators are draggable if you want more room for a column.

The top bar of the Accuracy Lab interface, showing the filter control, progress, and account information.

The top bar is deliberately minimal: how many rows the queue holds, the *Filter* control, your progress and your account. There is no site selector — the queue already spans the whole corpus, ordered as described in section 5.

Queue (left). Each row shows the site, the criterion and the technique (1.4.4 F94), preceded by the **xN** badge: how many findings of the exact same structural shape sit behind that case. The little triangle opens the list of siblings, if you want to see them one by one.

Filter is the only control: it shows all findings, only those you have not labelled yet, or only those you have already decided. It opens on the unlabelled ones.

Copy case (markdown)

▼ RIFERIMENTI WCAG

[1.4.4 Resize Text](#) Livello AA

[F94](#) tecnica · F94

▼ FINDING

key	snap_20260709_03b2e0:009359272476a1ff
code	WCAG2AA.Principle1.Guideline1_4.1_4_4.F94
type	warning (deterministic)
selector	#siti_tematici > div > div > div:nth-of-type(2) > div:nth-of-type(16) > div > a > div:nth-of-type(2) > h1

▼ MESSAGE

Element has font-size: calc(1.375rem + 1.5vw) (mixed viewport + relative units). Consider using only relative units.

▼ OUTERHTML

<h1 class="card-title h5">

The evidence panel. At the top, *WCAG references*: the criterion spelled out with its conformance level and the technique, both linked to the official W3C text — you do not have to decode the identifier to know what is being discussed. Below: the engine's message, the element's `outerHTML`, its attributes and the styles the browser *actually computed*. `type` distinguishes the nature of the finding: *error (deterministic)* is a violation the engine asserts with mechanical certainty — the kind of case where your judgement carries the most weight. *Copy case (markdown)* copies the whole dossier to the clipboard, handy if you want to discuss one in writing.

4. Three views of the same finding

The centre panel has three tabs, all on the same element. Use them in sequence: the first says *how the page is built*, the second *what a sighted user sees*, the third *what the code actually says*.

- **Rendered page** — the frozen page, live and navigable, with the reported element highlighted and scrolled into view automatically.
- **Screenshot** — the full-page capture with the element outlined. You can compare `original.png` (the live site at capture time) with `frozen.png` (the frozen copy): if they differ materially, the copy lost something and the case deserves caution.
- **HTML** — the full source, with the reported element highlighted among its neighbours: often that is where the case resolves itself. The *Scroll to finding* button takes you back to the element if you lose it while scrolling.

5. How to label

The verdict saves immediately. There is no confirmation step and the key does nothing. As a result, **the reason must be written BEFORE choosing the verdict**, or it will not be attached. The field says so itself: "*reason (optional) — write it BEFORE choosing*".

- 1 **Read the message** in the right-hand panel: it is the precise claim the engine is making. The WCAG references at the top tell you which criterion it is making it against.
- 2 **Check it on the element** using the three views. The question is not "is this page accessible?" but "is this specific claim, about this specific element, true?".
- 3 **Write the reason** in the field at the bottom right, whenever the case is not obvious. It is optional, but it is the part we need most: the verdict tells us *what*, the reason tells us *why*, and it is the reason that becomes a rule.
- 4 **Give the verdict** — for tp, for fp, for uncertain (or the three buttons at the bottom right). **It saves itself.**
- 5 **Move to the next** with (to go back).

/ next/prev · / / label (saves immediately; click another verdict to relabel)

The bar at the bottom of the window always repeats the commands and reminds you that the verdict is immediate.

Labels are **append-only**: if you come back to a case and change your mind, the new label supersedes the previous one and that one stays in the history. Nothing is deleted and there is no way to damage work already done — so do not hesitate to correct yourself.

The queue's order is chosen to make your hour count, and it changes only what reaches you first, never what is true:

- **It starts from the unlabelled**, because a session exists to cover ground the corpus does not yet cover.
- **It puts the hard cases first**, along with the largest patterns. Groups someone has already judged drop to the bottom.
- **It groups by shape**: a single representative for each structural pattern. The pattern is computed from the shape of the markup, not from the site: the same pattern legitimately spans different sites.

The xN badge is informative, not multiplicative. If you label a representative, you have labelled *that* finding and no other: propagating your judgement to the other N–1 would manufacture human judgements on cases no human ever opened, and that is exactly the corruption this lab exists to prevent. The number tells you how much the *family* weighs, so you can choose where to spend your attention.

6. When two experts disagree

If two **different** experts give opposite verdicts on the same finding, the system **does not choose**. The case becomes **contested**: not tp, not fp, and — crucially — **not the most recent verdict**. Recency plays no part anywhere in this path.

The reason is simple: between two competent people who looked at the same case, order of arrival is not an argument. Letting the last one win would produce a ground truth that depends on who opened the lab

second.

- **A disagreement is not an accident to be hidden:** it almost always marks a genuinely hard case, and that is valuable information. Do not try to guess what a colleague would have said — you cannot see it, and that is by design.
- **Coordination arbitrates.** An arbitration verdict settles the contest and the finding reads as tp or fp again, **but the disagreement stays on record:** an arbitration that erased what it arbitrated would be worse than no arbitration.
- **You need do nothing** when a case ends up contested: you are not notified and you are not asked to revisit it. Carry on working as normal.

Changing your mind about your *own* judgement is a different matter: that is self-correction, and there the most recent one wins (section 5).

7. Three practical rules

1. **The browser is authoritative, not the source.** If the HTML and Chromium's accessibility tree tell different stories, the tree is right. That is why the right-hand panel shows you the *computed* styles and not the declared ones.
2. **"Not decidable" is a good answer.** `u` plus one line of reasoning is worth far more than a guessed tp or fp: it tells us that case needs evidence we do not yet produce, and that is actionable.
3. **If the copy looks suspect, say so.** When you suspect the finding depends on something that did not load in the frozen copy — an image, a stylesheet — compare `original.png` with `frozen.png` in the Screenshot tab: if the copy lost the part that matters, label *uncertain* and write that down. It is not your mistake, it is a limit of our capture, and we want to know about it.

8. Quick reference

KEY	ACTION
<code>j</code> / <code>k</code>	next / previous finding
<code>t</code> / <code>f</code> / <code>u</code>	verdict tp / fp / uncertain — saves immediately
CONTROL	WHAT IT DOES
Filter	all · unlabelled only · labelled only
Reason field	the why behind the verdict — must be written before choosing
xN badge	how many findings of the same shape it stands for; the triangle opens them
Copy case (markdown)	copies the full case dossier to the clipboard
Scroll to finding	in the HTML tab, scrolls the source back to the reported element
Esci (log out)	ends the session